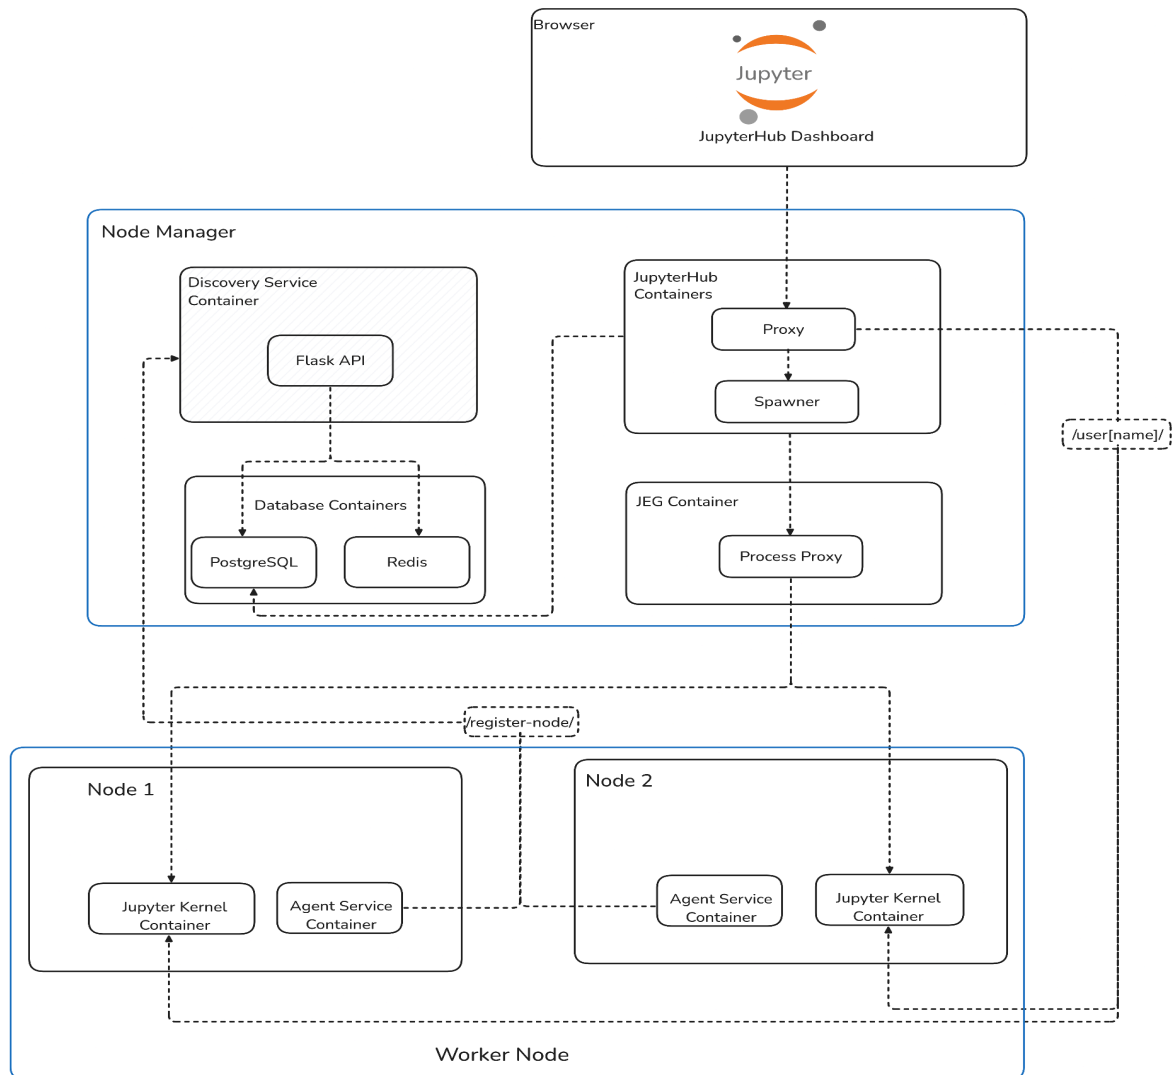


GUIDELINES

A. Arsitektur



B. Prasyarat

- Python3.10, 3.11
- Instal Docker dan Docker Compose di setiap node
- NVIDIA Driver & nvidia-container-toolkit (untuk akses GPU)
- 2 atau lebih node dalam jaringan lokal (1 control node, sisanya worker)
- Clone project: <https://github.com/danielcristho/jupyterhub-jeg-cluster.git> di control node, di worker opsional.
- Edit `/lib/systemd/system/docker.service` di setiap node untuk mengaktifkan akses Docker via jaringan (port 2375):
 - Cari baris **ExecStart**:

```
$ sudo nano /lib/systemd/system/docker.service
```

```
ExecStart=/usr/bin/dockerd -H fd:// -H tcp://0.0.0.0:2375
```
 - Reload dan restart Docker service:

```
sudo systemctl daemon-reload
```

```
sudo systemctl daemon-reexec
```

```
sudo systemctl restart docker.service
```

1. Konfigurasi Service Discovery

- a. Konfigurasi environment, ubah sesuai kebutuhan:

```
$ cd service-discovery
$ cp .env.example .env
$ cat .env
# Application Configuration
API_HOST=0.0.0.0
API_PORT=15002
FLASK_DEBUG=True
SECRET_KEY=secret-service111111

# Database Configuration
POSTGRES_HOST=postgres
POSTGRES_PORT=5432
POSTGRES_DB=voyager
POSTGRES_USER=postgres
POSTGRES_PASSWORD=postgres

REDIS_HOST=redis
REDIS_PASSWORD=redis@pass
REDIS_PORT=16379
REDIS_EXPIRE_SECONDS=45

SQLALCHEMY_TRACK_MODIFICATIONS=False
SQLALCHEMY_ECHO=False
```

- b. Ubah file init.sql (postgresql) dan redis.conf (redis) untuk mengatur username dan password:

```
-- DB Service Discovery
CREATE DATABASE voyager;
-- DB JupyterHub
CREATE DATABASE apollo;
GRANT ALL PRIVILEGES ON DATABASE voyager TO postgres;
GRANT ALL PRIVILEGES ON DATABASE apollo TO postgres;

CREATE EXTENSION IF NOT EXISTS "uuid-oss";
SET client_encoding = 'UTF8';
```

- c. Setelah menjalankan Docker Compose, pastikan terdapat tiga container yang running (discovery-api, redis, dan postgres):

```
$ docker compose up -d --build
```

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND NAMES	CREATED	STATUS	PORTS
68b27296d9da	service-discovery_discovery	"python app.py"	23 minutes ago	Up 23 minutes	
0.0.0.0:15002->15002/tcp, :::15002->15002/tcp		discovery-api			
9616551e7b2e	redis:7-alpine	"docker-entrypoint.s..."	23 minutes ago	Up 23 minutes (healthy)	6379/tcp,
0.0.0.0:16379->16379/tcp, :::16379->16379/tcp		redis			
d019c8b667ed	postgres:14-alpine	"docker-entrypoint.s..."	23 minutes ago	Up 23 minutes (healthy)	
0.0.0.0:5432->5432/tcp, :::5432->5432/tcp		postgres			

- d. Migrasi database dan upgrade skema, pastikan prosesnya berhasil:

```
$ docker exec -it discovery-api flask db init
$ docker exec -it discovery-api flask db migrate -m "START MIGRATION"
$ docker exec -it discovery-api flask db upgrade
```

- e. Akses API service discovery:

```
$ curl http://<CONTROL_NODE_IP>:15002/health-check
```

```
→ curl http://127.0.0.1:15002/health-check
{
  "database": {
    "postgres": "connected",
    "redis": "connected"
  },
  "message": "Hello, from [DiscoveryAPI]",
  "status": "ok"
}
```

- f. Akses postgresql dan redis:

```
$ docker exec -it postgres psql -U postgres
```

```
psql (14.18)
Type "help" for help.

postgres=# \l
               List of databases
  Name  | Owner  | Encoding | Collate  | Ctype    | Access privileges
-----+-----+-----+-----+-----+-----
 apollo | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =Tc/postgres +
        |          |          |          |          | postgres=CTc/postgres
 ds     | postgres | UTF8     | en_US.utf8 | en_US.utf8 |
 postgres | postgres | UTF8     | en_US.utf8 | en_US.utf8 |
 template0 | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =c/postgres +
        |          |          |          |          | postgres=CTc/postgres
 template1 | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =c/postgres +
        |          |          |          |          | postgres=CTc/postgres
 voyager | postgres | UTF8     | en_US.utf8 | en_US.utf8 | =Tc/postgres +
        |          |          |          |          | postgres=CTc/postgres
(6 rows)

postgres=#
```

TablePlus 1.2.2 (258) (Trial)

Discovery Service

Discovery Service

PostgreSQL NO TAG

127.0.0.1 5432

postgres postgres

voyager

PREFERRED

Key... Cert... CA Cert... -

☐ SSH Tunnel

Save Test Connect

PostgreSQL 14.18 : Discovery

id	hostname	ip	cpu_cores	ram_gb	has_gpu	gpu_info	is_active
1	chrstdan	172.18.0...	4	18.79	FALSE	[]	TRUE

Items Queries History

Search for items...

Table

- alembic_version
- node_metrics
- node_selections
- nodes
- profiles

```
$ docker exec -it redis redis-cli -a redis@pass -p 16379
```

```
+ docker exec -it redis redis-cli -a redis@pass -p 16379
Warning: Using a password with '-a' or '-u' option on the command line interface may not be safe.
127.0.0.1:16379> keys *
1) "node:chrstdan:info"
2) "node:chrstdan:ip"
127.0.0.1:16379> get node:chrstdan:info
"{\"hostname\": \"chrstdan\", \"ip\": \"172.18.0.1\", \"cpu_cores\": 4, \"gpu_info\": [], \"has_gpu\": false, \"disk_usage_percent\": 72.6, \"disk_usage_percent\": 83.5, \"active_jupyterlab\": 0, \"active_ray\": 0, \"active_ray\": 0}"
127.0.0.1:16379>
```

2. Konfigurasi Service Agent

- a. Akses node worker, gunakan command docker run untuk menjalankan service agent:

```
# Menjalankan agent di setiap node
```

```
docker run --name agent -d \
--net=host \
-e DISCOVERY_URL=http://<CONTROL_NODE_IP>:15002/register-node \
-v /var/run/docker.sock:/var/run/docker.sock \
danielcristh0/agent:1.1
```

```
# Tambahkan "--gpus all" di node yang memiliki GPU
```

```
docker run --name agent -d \
--net=host \
-e DISCOVERY_URL=http://<CONTROL_NODE_IP>:15002/register-node \
-v /var/run/docker.sock:/var/run/docker.sock \
--gpus all \
danielcristh0/agent:1.1
```

```
[DEBUG] Container Summary: Total=55, JupyterLab=2, Ray=7
[DEBUG] Container Summary: Total=55, JupyterLab=2, Ray=7
[DEBUG] Send Info: {'hostname': 'rpl', 'ip': '10.21.73.122', 'cpu': 24, 'gpu':
[{'name': 'NVIDIA GeForce RTX 3080 Ti', 'index': 0, 'uuid': 'GPU-56c3e796-124
e-f059-16a8-f9be2b254ce0', 'memory_total_mb': 12288, 'memory_used_mb': 7632, '
memory_util_percent': 62.11, 'utilization_gpu_percent': 0, 'temperature_gpu':
56}], 'has_gpu': True, 'ram_gb': 67.11, 'cpu_usage_percent': 2.5, 'memory_usag
e_percent': 16.1, 'disk_usage_percent': 42.3, 'active_jupyterlab': 2, 'active
ray': 7, 'total_containers': 55, 'last_updated': '2025-06-19T20:45:16.193296Z'
}
[AGENT] Registered: rpl (10.21.73.122) → 200
□
```

- b. Cek database service discovery untuk memastikan apakah node telah berhasil didaftarkan:

PostgreSQL 14.18 : Discovery Service : voyager : public.nodes (Trial)

id	hostname	ip	cpu_cores	ram_gb	has_gpu	gpu_info	is_active	max_container	created_at	updated_at
201	rpl	10.21.73.139	24	67.11	FALSE	[]	TRUE	10	2025-07-08 20:34:14.2...	2025-07-10 10:23:13.4...
25	rpl-02	10.21.73.125	20	33.38	FALSE	[]	TRUE	10	2025-07-08 16:59:46.5...	2025-07-10 10:23:17.1...
3	chrstdan	10.125.180...	4	18.79	FALSE	[]	TRUE	10	2025-06-22 00:32:27.1...	2025-07-10 10:23:19.4...
1	worker1	192.168.122...	2	4.1	FALSE	[]	FALSE	10	2025-06-21 14:21:28.1...	2025-07-08 18:10:43.7...
35	rpl-1	10.21.73.107	20	16.5	FALSE	[]	FALSE	10	2025-07-08 19:52:57.4...	2025-07-08 20:50:04.8...
2	worker2	192.168.122...	2	4.1	FALSE	[]	FALSE	10	2025-06-21 14:21:28.1...	2025-07-08 16:36:37.1...

- c. Cek node yang tersedia di service discovery, dari endpoint `available-nodes`:
- ```
curl http://127.0.0.1:15002/available-nodes
```

```
{
 "all_available_nodes": [
 {
 "active_jupyterlab": 0,
 "active_ray": 0,
 "cpu_cores": 4,
 "cpu_usage_percent": 11.8,
 "created_at": "2025-08-06T09:12:07.012743",
 "disk_usage_percent": 83.5,
 "gpu_info": [],
 "has_gpu": false,
 "hostname": "chrstdan",
 "id": 1,
 "ip": "172.18.0.1",
 "is_active": true,
 "load_score": 68.32,
 "max_containers": 10,
 "memory_usage_percent": 73.6,
 "ram_gb": 18.79,
 "total_containers": 0,
 "updated_at": "2025-08-06T10:59:44.964723"
 }
],
 "load_balancing": {
 "algorithm": "round_robin",
 "requested_count": 1,
 "round_robin_counter": 3,
 "selected_count": 1
 },
 "selected_nodes": [
 {
 "active_jupyterlab": 0,
 "active_ray": 0,
 "cpu_cores": 4,
 "cpu_usage_percent": 11.8,
 "created_at": "2025-08-06T09:12:07.012743",
 "disk_usage_percent": 83.5,
 "gpu_info": [],
 "has_gpu": false,
 "hostname": "chrstdan",
 "id": 1,
 "ip": "172.18.0.1",
 "is_active": true,
 "load_score": 68.32,
 "max_containers": 10,
 "memory_usage_percent": 73.6,
 "ram_gb": 18.79,
 "total_containers": 0,
 "updated_at": "2025-08-06T10:59:44.964723"
 }
],
 "total_available_nodes": 1
}
```

### 3. Konfigurasi Jupyter

#### a. Konfigurasi Jupyter Enterprise Gateway

- Konfigurasi SSH Passwordless:  

```
$ ssh-keygen -t rsa -b 4096 -N "" -f ./jeg/ssh_keys/id_rsa
```
- Tambahkan public key ke node target, ulangi langkah ini di setiap node yang ada:  

```
$ ssh-copy-id -i ~/.ssh/id_rsa.pub user@<ip_node_worker>
```
- Agar JEG tidak gagal karena verifikasi host SSH, tambahkan fingerprint dari node ke `known_hosts`:  

```
$ ssh-keyscan -H <ip_node_worker> >> ./jeg/ssh_keys/known_hosts
```
- Verifikasi, jika tidak diminta password, maka konfigurasi berhasil:  

```
$ ssh user@<ip_node_worker>
```

#### b. Konfigurasi environment, ubah sesuai kebutuhan:

```
$ cd jupyterlab
$ cp .env.example .env
```

```

$ cat .env
$ cd jupyterlab
JupyterHub Admin
JUPYTERHUB_ADMIN=admin

Docker Image Default (optional)
DOCKER_NOTEBOOK_IMAGE=danielcrith0/jupyterlab:cpu

Proxy Auth Token
CONFIGPROXY_AUTH_TOKEN=abc123

Database Config
POSTGRES_HOST=<IP_CONTROL_NODE>
POSTGRES_PORT=5432
POSTGRES_DB=apollo
POSTGRES_USER=postgres
POSTGRES_PASSWORD=postgres

JEG Config
EG_LOG_LEVEL=DEBUG
EG_AUTH_TOKEN=jeg-jeg-an
EG_REMOTE_USER=<SSH_REMOTE_USER>
EG_RESPONSE_IP=<IP_CONTROL_NODE>

Grafana
GRAFANA_ADMIN_PASSWORD=monitoring

```

- c. Update nilai dari **API\_URL** di file `jupyterlab/hub/form/main.js` dengan IP control node. Tujuannya adalah untuk mengintegrasikan JupyterHub dengan Service Discovery:
 

```
const API_URL = "http://<IP_CONTROL_NODE>:15002";
```
- d. Pull Docker image JupyterLab di setiap node worker untuk mempercepat proses spawn di JupyterHub. Tujuannya agar image sudah tersedia, sehingga proses spawn container dapat berjalan lebih cepat dan tanpa delay.:
  - Semua node wajib pull image CPU:
 

```
$ docker pull danielcrith0/jupyterlab:cpu
```
  - Jika node memiliki GPU, pull juga image GPU:
 

```
$ docker pull danielcrith0/jupyterlab:gpu
```
- e. Jalankan Docker Compose Jupyter, pastikan ada 3 container yang running (jeg, proxy & jupyterhub):
 

```
$ cd jupyterlab
$ docker compose up -d --build
```

```

NAMES
57a234355450 danielcrith0/gateway:1.1 "/usr/local/bin/entr..." 16 minutes ago Up 16 minutes (unhealthy)
:::8889->8889/tcp, 8888/tcp jeg
b14874d8a8b6 jupyterhub/configurable-http-proxy:4 "/srv/configurable-h..." 16 minutes ago Up 16 minutes
proxy
300a0c55c035 prom/prometheus:latest "/bin/prometheus --c..." 16 minutes ago Up 16 minutes
prometheus
cfa172022867 danielcrith0/jupyterhub:1.1 "tini -- jupyterhub ..." 16 minutes ago Up 16 minutes (unhealthy)
jupyterhub

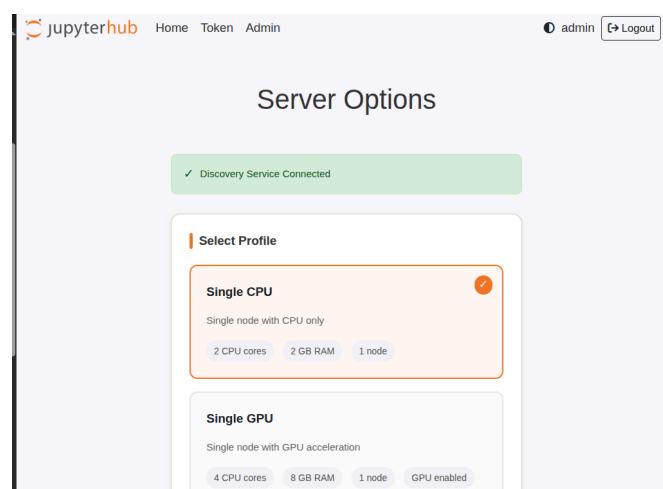
```

f. Akses halaman JupyterHub di `http://<IP_CONTROL_NODE>:18000`:

- Lakukan sign up terlebih dahulu. Untuk akun admin, gunakan username: admin dan password bebas (minimal 8 karakter). Setelah sign up sebagai admin, akun akan otomatis memiliki hak akses admin
- Jika sudah, lalu login kemudian klik “Start My Server” untuk memilih node yang digunakan untuk menjalankan JupyterLab.



The image shows the JupyterHub Sign In page. At the top, there is a JupyterHub logo and a 'Login' button. Below this is a 'Sign In' section with a warning message: 'Warning: JupyterHub seems to be served over an unsecured HTTP connection. We strongly recommend enabling HTTPS for JupyterHub.' Below the warning, there are input fields for 'Username:' and 'Password:', followed by a 'Sign In' button. At the bottom of the sign in section, there is a link that says 'Sign up to create a new user.'



The image shows the JupyterHub Server Options page. At the top, there is a JupyterHub logo and navigation links for 'Home', 'Token', and 'Admin'. On the right, there is a user profile 'admin' and a 'Logout' button. The main heading is 'Server Options'. Below this, there is a green status bar that says 'Discovery Service Connected'. Underneath, there is a 'Select Profile' section. It contains two profile options: 'Single CPU' (which is selected, indicated by an orange border and a checkmark) and 'Single GPU'. The 'Single CPU' profile is described as 'Single node with CPU only' and has '2 CPU cores', '2 GB RAM', and '1 node'. The 'Single GPU' profile is described as 'Single node with GPU acceleration' and has '4 CPU cores', '8 GB RAM', '1 node', and 'GPU enabled'.



Environment Configuration

Docker Image

CPU Environment (danielcris0/jupyterlab.cpu) ▾

Choose CPU for general computing or GPU for machine learning tasks

Node Configuration

Single Node

Available Nodes

chrstdan (Primary) • Moderate Load

IP: 172.18.0.1

CPU: 4 cores

Memory: 18.79 GB

Containers: 0 active

CPU USAGE

MEMORY USAGE

37.4%

78.0%

Selection Summary

Selected 1 node:

• Total Resources: 4 CPU cores, 18.79 GB RAM

Start

Available Nodes

rpl (Primary) • Available

IP: 10.21.73.122

CPU: undefined cores

Memory: 67.11 GB

Containers: 55 active

GPU: GPU Available

CPU USAGE

MEMORY USAGE

2.0%

16.1%

Selection Summary

Selected 1 node:

• Total Resources: 0 CPU cores, 67.11 GB RAM, 1 GPU

• Average Load: CPU 2.0%, Memory 16.1%

Start

Your server is starting up.

You will be redirected automatically when it's ready for you.

Server ready at </user/admin/>

#### ▼ Event log

Server requested

Spawning server...

Server ready at </user/admin/>

The screenshot displays the JupyterLab web interface. On the left is a file browser sidebar with a search bar and a table of files. The main area is divided into three horizontal panes: Notebook, Console, and Other. The Notebook pane shows two kernels: 'Python 3 (ipykernel)' and 'Python 3 on rpl'. The Console pane also shows the same two kernels. The Other pane contains icons for Terminal, Text File, Markdown File, Python File, and Show Contextual Help. A small notification box in the bottom right corner asks if the user wants to receive official Jupyter news.

| Name                  | Last Modified |
|-----------------------|---------------|
| kaggle                | 15 hours ago  |
| cifar10_image_proc... | 8 hours ago   |
| rf_image_processin... | 13 hours ago  |
| Untitled.ipynb        | a minute ago  |
| Untitled2.ipynb       | 8 hours ago   |

Would you like to receive official Jupyter news?  
Please read the privacy policy.

Open privacy policy Yes No